

Importing the libraries

```
In [95]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
import math as m
from sklearn import linear_model
import warnings
warnings.filterwarnings("ignore")
```

Splitting data into train and test

```
In [106]: X, y = make_classification(n_samples=50000, n_features=15, n_informativ
e=10, n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7, r
andom_state=15)
```

```
In [107]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
5, random_state=15)
```

```
In [108]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[108]: ((37500, 15), (37500,), (12500, 15), (12500,))
```

Intializing initial parameters

```
In [109]: w = np.zeros_like(X_train[0])# initial weight vector
b = 0      # initial intercept value
eta0 = 0.0001 # learning rate
```

```
alpha = 0.0001 # lambda value
N = len(X_train)
```

```
In [110]: def sigmoid(w,x,b):
          return 1/(1+np.exp(-(np.dot(x,w.T)+b))) #return 1/1+e(-x)
```

```
In [111]: def logloss(w,x,y,b,reg=0):
          val=sigmoid(w,x,b)
          return -np.mean(y*np.log10(val)+(1-y)*np.log10(1-val))+reg # cost
          function of logistic regression
```

```
In [112]: print("INITIAL LOG LOSS:")
          logloss(w,X_train,y_train,b)
```

INITIAL LOG LOSS:

Out[112]: 0.3010299956639812

SGD Alorithm

```
In [113]: def sgd_algo(x_train,y_train,x_test,y_test,eta0,alpha,w,b,epoch):
          train_loss=[]
          test_loss=[]
          epoc=[]
          for i in range(0,epoch):
              epoc.append(i)
              for j in range(0,N):
                  reg=alpha/2*np.dot(w.T,w) #regulrization term
                  w = ((1-eta0*(alpha/N))*w)+((eta0*x_train[j])*(y_train[j]-s
                  igmoid(w,x_train[j],b))) # updating weight vector
                  b = b+(eta0*(y_train[j]-sigmoid(w,x_train[j],b))) # u
                  pdatind intercept
                  train=logloss(w,x_train,y_train,b,reg)
                  train_loss.append(train) # calculating train and test loss f
                  or updated w,b on each epoch
                  test=logloss(w,x_test,y_test,b,reg)
                  test_loss.append(test)
```

```

        """
        if i==0 :
            continue                                # blo
        ck to check convergence
        else:
            checking not getting optimum value as sklearn implementation      # but
            if abs(train_loss[i]-train_loss[i-1])>.001:
                continue
            else:
                break"""
    return w,b ,train_loss,test_loss,epoc

```

```

In [114]: epoch=6
w,b,tr,te,epoc=sgd_algo(X_train,y_train,X_test,y_test,eta0, alpha,w,b,e
poch)

```

```

In [115]: print("optimal weight vector:")
print("\n")
print(w)

```

optimal weight vector:

```

[-0.40819512  0.18608803 -0.13873193  0.33720386 -0.19058573  0.5532163
1
-0.44625706 -0.09550674  0.20929646  0.16027648  0.18715532  0.0082206
4
-0.0701101   0.33795986  0.0206259 ]

```

```

In [116]: print("optimal intercept value:")
print("\n")
print(b)

```

optimal intercept value:


```

-0.7613483856830058

```

```
In [120]: %matplotlib inline
import matplotlib.pyplot as plt
plt.figure(figsize=(8,6))
plt.grid()
plt.plot(epoc,tr, label='tain log loss')
plt.plot(epoc,te, label='test log loss')
plt.scatter(epoc,tr)
plt.scatter(epoc,te)
plt.title('log loss curve of logistic regression')
plt.legend()
```

Out[120]: <matplotlib.legend.Legend at 0x22f1df62f60>



sklearn implementation

```
In [97]: clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log',
```

```
random_state=15, penalty='l2', tol=.001, verbose=2, learning_rate='constant')
clf.fit(X=X_train, y=y_train)
```

```
-- Epoch 1
Norm: 0.76, NNZs: 15, Bias: -0.314605, T: 37500, Avg. loss: 0.455801
Total training time: 0.02 seconds.
-- Epoch 2
Norm: 0.92, NNZs: 15, Bias: -0.469578, T: 75000, Avg. loss: 0.394737
Total training time: 0.04 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580452, T: 112500, Avg. loss: 0.385561
Total training time: 0.06 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.660824, T: 150000, Avg. loss: 0.382161
Total training time: 0.08 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.717218, T: 187500, Avg. loss: 0.380474
Total training time: 0.11 seconds.
-- Epoch 6
Norm: 1.06, NNZs: 15, Bias: -0.761816, T: 225000, Avg. loss: 0.379481
Total training time: 0.13 seconds.
Convergence after 6 epochs took 0.13 seconds
```

```
Out[97]: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
eta0=0.0001, fit_intercept=True, l1_ratio=0.15,
learning_rate='constant', loss='log', max_iter=None, n_iter=None,
e,
n_jobs=1, penalty='l2', power_t=0.5, random_state=15, shuffle=True,
ue,
tol=0.001, verbose=2, warm_start=False)
```

```
In [100]: clf.coef_
```

```
Out[100]: array([[ -0.41177431,  0.18416782, -0.13895073,  0.33572511, -0.1842323
7,
0.5494352 , -0.45213692, -0.08857465,  0.21536661,  0.1735175
7,
```

```
0.18480827, 0.00443463, -0.07033001, 0.33683181, 0.0200412  
9]])
```

```
In [99]: clf.coef_-w
```

```
Out[99]: array([[ -0.0035792 , -0.00192022, -0.0002188 , -0.00147875,  0.0063533  
6,          -0.0037811 , -0.00587986,  0.00693208,  0.00607015,  0.0132411  
,          -0.00234705, -0.00378601, -0.00021991, -0.00112805, -0.0005846  
2]])
```

- difference between sklearn and custom implementation is almost in terms of 10^{-3}

```
In [123]: def pred(w,b, X):  
    N = len(X)  
    predict = []  
    for i in range(N):  
        if sigmoid(w, X[i], b) >= 0.5:  
            predict.append(1)  
        else:  
            predict.append(0)  
    return np.array(predict)  
print("train accuracy :")  
print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))  
print("test accuracy :")  
print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))  
  
train accuracy :  
0.9617866666666667  
test accuracy :  
0.95952
```