

8_D

Importing libraries

```
In [12]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import GridSearchCV
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

Reading the dataset

```
In [13]: data = pd.read_csv('task_d.csv')
```

```
In [14]: data.head()
```

Out[14]:

	x	y	z	x*x	2*y	2*z+3*x*x	w	target
0	-0.581066	0.841837	-1.012978	-0.604025	0.841837	-0.665927	-0.536277	0
1	-0.894309	-0.207835	-1.012978	-0.883052	-0.207835	-0.917054	-0.522364	0
2	-1.207552	0.212034	-1.082312	-1.150918	0.212034	-1.166507	0.205738	0

	x	y	z	x*x	2*y	2*z+3*x*x	w	target
3	-1.364174	0.002099	-0.943643	-1.280666	0.002099	-1.266540	-0.665720	0
4	-0.737687	1.051772	-1.012978	-0.744934	1.051772	-0.792746	-0.735054	0

```
In [15]: X = data.drop(['target'], axis=1).values
Y = data['target'].values
```

1. Finding the Correlation between the features

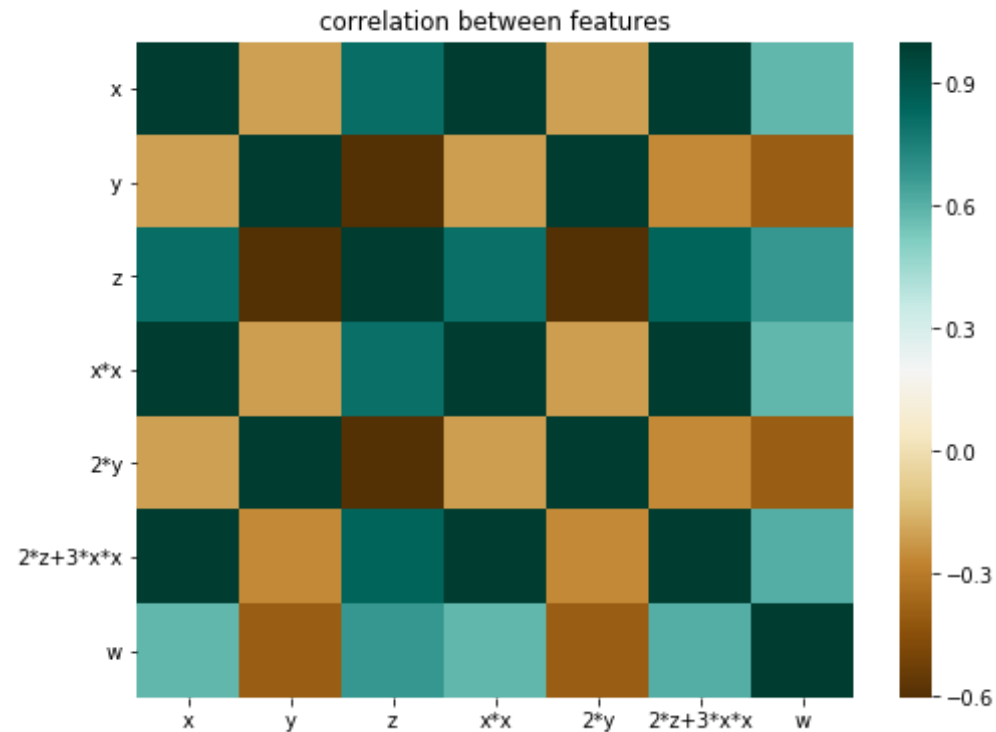
```
In [16]: corr=data[data.columns[:-1]].corr()
corr
```

Out[16]:

	x	y	z	x*x	2*y	2*z+3*x*x	w
x	1.000000	-0.205926	0.812458	0.997947	-0.205926	0.996252	0.583277
y	-0.205926	1.000000	-0.602663	-0.209289	1.000000	-0.261123	-0.401790
z	0.812458	-0.602663	1.000000	0.807137	-0.602663	0.847163	0.674486
x*x	0.997947	-0.209289	0.807137	1.000000	-0.209289	0.997457	0.583803
2*y	-0.205926	1.000000	-0.602663	-0.209289	1.000000	-0.261123	-0.401790
2*z+3*x*x	0.996252	-0.261123	0.847163	0.997457	-0.261123	1.000000	0.606860
w	0.583277	-0.401790	0.674486	0.583803	-0.401790	0.606860	1.000000

```
In [17]: plt.figure(figsize=(8,6))
plt.title("correlation between features")
sns.heatmap(corr, cmap="BrBG",
            xticklabels=corr.columns,
            yticklabels=corr.columns)
```

Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x24c2f177470>



Task: 1 Logistic Regression

```
In [18]: alpha = np.logspace(-5, 8, 10)
print(alpha)
param_grid={'C':alpha}
logreg = LogisticRegression()

[1.00000000e-05 2.78255940e-04 7.74263683e-03 2.15443469e-01
 5.99484250e+00 1.66810054e+02 4.64158883e+03 1.29154967e+05
 3.59381366e+06 1.00000000e+08]
```

2. Finding the best model for the given data

```
In [19]: logreg = GridSearchCV(logreg, param_grid, cv=5)
```

```
In [20]: logreg.fit(X,Y)
```

```
Out[20]: GridSearchCV(cv=5, error_score='raise',  
                    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,  
                    fit_intercept=True,  
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=  
                    1,  
                    penalty='l2', random_state=None, solver='liblinear', tol=0.00  
                    01,  
                    verbose=0, warm_start=False),  
                    fit_params=None, iid=True, n_jobs=1,  
                    param_grid={'C': array([1.00000e-05, 2.78256e-04, 7.74264e-03,  
                    2.15443e-01, 5.99484e+00,  
                    1.66810e+02, 4.64159e+03, 1.29155e+05, 3.59381e+06, 1.00000e+0  
                    8])},  
                    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',  
                    scoring=None, verbose=0)
```

```
In [21]: logreg.best_params_
```

```
Out[21]: {'C': 1e-05}
```

3. Getting the weights with the original data

```
In [22]: best_model=LogisticRegression(C=1e-05)
```

```
In [23]: best_model.fit(X,Y)
```

```
Out[23]: LogisticRegression(C=1e-05, class_weight=None, dual=False, fit_intercep  
t=True,  
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=  
                    1,  
                    penalty='l2', random_state=None, solver='liblinear', tol=0.00  
                    01,  
                    verbose=0, warm_start=False)
```

```
In [24]: predictions = best_model.predict(X)
```

```
In [25]: accu=accuracy_score(Y, predictions)
print(accuracy_score(Y, predictions))
```

```
1.0
```

```
In [26]: wei=best_model.coef_[0]
print(best_model.coef_)
```

```
[[ 0.00036369 -0.000345    0.00048449  0.00035933 -0.000345    0.000381
89
  0.00032048]]
```

4. Modifying original data

```
In [29]: X_NEW=X+.01 # ADDING NOISE
```

```
In [30]: updated_model=best_model.fit(X_NEW,Y)
```

```
In [31]: prediction = best_model.predict(X_NEW)
```

```
In [32]: new_accu=accuracy_score(Y, prediction)
print(accuracy_score(Y, prediction))
```

```
1.0
```

```
In [33]: w_new=updated_model.coef_
w_new
```

```
Out[33]: array([[ 0.00036369, -0.000345 ,  0.00048449,  0.00035933, -0.000345
,
  0.00038189,  0.00032048]])
```

5. Checking deviations in metric and weights

```
In [34]: print(new_accu-accu)
```

```
0.0
```

```
In [35]: difference=abs((wei-w_new))[0]
print(difference)
```

```
[3.04620171e-11 3.04801046e-11 3.04930581e-11 3.04600787e-11
 3.04801046e-11 3.04634312e-11 3.04687841e-11]
```

```
In [36]: n=len(data.columns)-1
percentage_change=[]
for i in range(n):                                # calculating the percentage change in weight
    cp=(difference[i]/wei[i])*100
    percentage_change.append(cp)
```

```
In [37]: columns=list(data.columns.values)
indices=sorted(range(len(percentage_change)), key=lambda i: percentage_change[i])[-4:]
print("the top 4 features which have higher % change in weights ")
for j in indices:
    print(columns[j])
```

```
the top 4 features which have higher % change in weights
2*z+3*x*x
x
x*x
w
```

Task: 2 Linear SVM

2. Finding the best model for the given data

```
In [38]: alpha = np.logspace(-5, 8, 10)
print(alpha)
param_grid={'C':alpha}
svm = SVC(kernel="linear")

[1.00000000e-05 2.78255940e-04 7.74263683e-03 2.15443469e-01
 5.99484250e+00 1.66810054e+02 4.64158883e+03 1.29154967e+05
 3.59381366e+06 1.00000000e+08]
```

```
In [39]: model = GridSearchCV(svm, param_grid, cv=5)
```

```
In [40]: model.fit(X,Y)
```

```
Out[40]: GridSearchCV(cv=5, error_score='raise',
      estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.
0,
      decision_function_shape='ovr', degree=3, gamma='auto', kernel='linea
r',
      max_iter=-1, probability=False, random_state=None, shrinking=True,
      tol=0.001, verbose=False),
      fit_params=None, iid=True, n_jobs=1,
      param_grid={'C': array([1.00000e-05, 2.78256e-04, 7.74264e-03,
2.15443e-01, 5.99484e+00,
      1.66810e+02, 4.64159e+03, 1.29155e+05, 3.59381e+06, 1.00000e+0
8])}),
      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
      scoring=None, verbose=0)
```

```
In [41]: model.best_params_
```

```
Out[41]: {'C': 0.007742636826811269}
```

3. Getting the weights with the original data

```
In [42]: best_model=SVC(kernel='linear',C=0.007742636826811269)
```

```
In [43]: best_model.fit(X,Y)
```

```
Out[43]: SVC(C=0.007742636826811269, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

```
In [44]: predictions = best_model.predict(X)
         accu=accuracy_score(Y, predictions)
         print(accuracy_score(Y, predictions))
```

```
1.0
```

```
In [45]: wei=best_model.coef_[0]
         print(best_model.coef_)
```

```
[[ 0.16056222 -0.20788705  0.32826166  0.14998082 -0.20788705  0.174615
 87
   0.13401176]]
```

4. Modifying original data

```
In [46]: X_NEW=X+.01
```

```
In [47]: updated_model=best_model.fit(X_NEW,Y)
```

```
In [48]: prediction = best_model.predict(X_NEW)
```

```
In [49]: new_accu=accuracy_score(Y, prediction)
         print(accuracy_score(Y, prediction))
```

```
1.0
```

```
In [50]: w_new=updated_model.coef_
```



```
In [51]: w_new
```

```
Out[51]: array([[ 0.16049457, -0.20810298,  0.32832289,  0.14997999, -0.20810298,
                0.17462251,  0.13395324]])
```

5. Checking deviations in metric and weights

```
In [52]: print(new_accu-accu)
```

```
0.0
```

```
In [53]: difference=(k-w_new)[0]
         print(difference)
```

```
-----
----
NameError                                Traceback (most recent call l
ast)
<ipython-input-53-633cf4288d5c> in <module>()
----> 1 difference=(k-w_new)[0]
      2 print(difference)

NameError: name 'k' is not defined
```

```
In [ ]: n=len(data.columns)-1
        percentage_change=[]
        for i in range(n):                # calculating the percentage chang
e in weight
            cp=(difference[i]/wei[i])*100
            percentage_change.append(cp)
```

```
In [54]: columns=list(data.columns.values)
         indices=sorted(range(len(percentage_change)), key=lambda i: percentage_
change[i])[-4:]
         print("the top 4 features which have higher % change in weights ")
```

```
for j in indices:  
    print(columns[j])
```

the top 4 features which have higher % change in weights

$2z + 3x \cdot x$

x

$x \cdot x$

w

Observations

- FROM CORRELATION MATRIX WE CAN SEE THAT SOME FEATURES ARE HIGHLY CORRELATED.
- AFTER DONE THE PERTURBATION TEST ON BOTH SVM AND LR, WEIGHT HAVE ONLY SMALL (VERY SMALL) CHANGE.
- SINCE NO ANY DRASTIC CHANGE WEIGHT WE CAN CONCLUDE THAT THERE IS NO COLLINEARITY BETWEEN FEATURES BASED ON PERTURBATION TEST.
- THE TOP 4 FEATURES WHICH HAVE HIGHER % CHANGE IN WEIGHTS ARE FEATURES WHICH HIGHLY CORRELATED WITH OTHER FEATURES.
(UNDERSTAND FROM CORRELATION MATRIX)